

Navigation

October 4, 2019

```
In [1]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 2.0.
```

Note: This code is made using help from various sources like udacity and various people's implementation. I'll try to improve upon this code as i learn

0.1 Deep Q-Network Architecture

The network is something like this:

1. input_size = state_size = 37
2. 2 linear layers, one with 128 nodes as output and other with 64 nodes as output with relu activation for both
3. A third linear layer with 4 outputs for 4 actions

0.2 Learning algo to use

The algo used in this project will be Q-learning algo with experience relay and fixed Q targets. The optimization of network will be done through Adam optimizer

0.2.1 Hyperparameters used

- BUFFER_SIZE = int(1e5) : replay buffer size
- BATCH_SIZE = 64 : minibatch size
- GAMMA = 0.99 : discount factor
- TAU = 1e-3 : for soft update of target parameters
- LR = 5e-4 : learning rate
- UPDATE_EVERY = 4 : how often to update the network
- EPS_DECAY=0.995 : the reduction factor of the epsilon-greedy policy

0.3 Environment Information

The environment we will use is [Unity MLAgents](#)

```
In [2]: from unityagents import UnityEnvironment
        from dqn_agent import Agent
        from collections import deque

        import numpy as np
        import random
        import os
        import torch
        import matplotlib.pyplot as plt
```

Next, we'll start the environment. To run on a new machine, before running the environment, change the `file_name` parameter to the path where the Unity environment is installed in the machine.

Here are the various paths for different operating systems:

- Mac: "path/to/Banana.app"
- Windows (x86): "path/to/Banana_Windows_x86/Banana.exe"
- Windows (x86_64): "path/to/Banana_Windows_x86_64/Banana.exe"
- Linux (x86): "path/to/Banana_Linux/Banana.x86"
- Linux (x86_64): "path/to/Banana_Linux/Banana.x86_64"
- Linux (x86, headless): "path/to/Banana_Linux_NoVis/Banana.x86"
- Linux (x86_64, headless): "path/to/Banana_Linux_NoVis/Banana.x86_64"

The environment has 4 actions: 0: walk forward 1: walk backward 2: turn left 3: turn right

State space has 37 dimensions. A reward of +1 will be awarded for collecting yellow banana and -1 for purple banana

```
In [3]: env = UnityEnvironment(file_name='/data/Banana_Linux_NoVis/Banana.x86_64')
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
    Number of Brains: 1
```

```
    Number of External Brains : 1
```

```
    Lesson number : 0
```

```
    Reset Parameters :
```

```
Unity brain name: BananaBrain
```

```
    Number of Visual Observations (per agent): 0
```

```
    Vector Observation space type: continuous
```

```
    Vector Observation space size (per agent): 37
```

```
    Number of stacked Vector Observation: 1
```

```
    Vector Action space type: discrete
```

```
    Vector Action space size (per agent): 4
```

```
    Vector Action descriptions: , , ,
```

```
In [4]: brain_name = env.brain_names[0]
        brain = env.brains[brain_name]
```

0.4 Train the agent with DQN

```
In [5]: agent = Agent(state_size=37, action_size=4, seed=32, fc1_units=128, fc2_units=64)
```

```
In [6]: save_path = 'checkpoint.pth'
        episode_num = 1
        def save_model(model, episode_num):
            if model is not None:
                checkpoint = {
                    'episode_num': episode_num,
                    'state_dict': model.state_dict()
                }
                torch.save(checkpoint, save_path)
                print('\r\tsave episode' + str(episode_num), end='')

        def load_model(model):
            if os.path.exists(save_path) and model is not None:
                checkpoint = torch.load(save_path)
                episode_num = checkpoint['episode_num']
                model.load_state_dict(checkpoint['state_dict'])
                print('Loaded episode number: ', episode_num)
                return episode_num
            return 1

        episode_num = load_model(agent.qnetwork_local)
        print(episode_num)
```

```
Loaded episode number: 700
700
```

```
In [7]: def dqnn(n_episodes=2000, max_t=100000, eps_start=1., eps_end=.01, eps_decay=.995):
        scores = []
        scores_window = deque(maxlen=100)
        eps = eps_start

        for i_episode in range(n_episodes):
            env_info = env.reset(train_mode=True)[brain_name]
            state = env_info.vector_observations[0]
            score = 0
            for t in range(max_t):
                action = (int)(agent.act(state, eps))
                env_info = env.step(action)[brain_name]
                next_state = env_info.vector_observations[0]
                reward = env_info.rewards[0]
                done = env_info.local_done[0]
                agent.step(state, action, reward, next_state, done)
                state = next_state
                score+=reward
```

```

        if done: break
    scores_window.append(score)
    scores.append(score)
    eps = max(eps_end, eps_decay * eps)
    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    if i_episode % 100 == 0:
        print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
        save_model(agent.qnetwork_local, i_episode)
    if np.mean(scores_window) >= 15.:
        print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))

        torch.save(agent.qnetwork_local.state_dict(), 'model.pth')
        break
    return scores

```


```
In [9]: scores = dqn()
```

```
Episode 785          Average Score: 15.05
```

```
Environment solved in 685 episodes!          Average Score: 15.05
```

```
In [10]: fig = plt.figure(figsize=(13, 10))
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episodes')
plt.show()

```



0.5 Testing the trained

```
In [8]: agent.qnetwork_local.load_state_dict(torch.load('model.pth'))
```

```
env_info = env.reset(train_mode=False)[brain_name]
```

```
state = env_info.vector_observations[0]
```

```
score = 0
```

```
while True:
```

```
    action = agent.act(state)
```

```
    env_info = env.step(action)[brain_name]
```

```
    next_state = env_info.vector_observations[0]
```

```
    reward = env_info.rewards[0]
```

```
    done = env_info.local_done[0]
```

```
    score+=reward
```

```
    state = next_state
```

```
    if done: break
```

```
print('Score:', score)
```

Score: 16.0

```
In [ ]: env.close()
```

0.6 Future work to consider:

I'm planning to add following features in this: * Duelling DQN * Double DQN * Prioritized Experienced Replay

```
In [ ]:
```